S&DS 365 / 665 Intermediate Machine Learning

## **Variational Inference and VAEs**

October 12

#### Reminders

- Practice midterm and solutions posted
- Multiple review sessions
- Exam Monday: Cheat sheet, parallels practice

## **For Today**

- Variational inference: The ELBO (rehash)
- Variational autoencoders
- Demo notebook
- Questions for review

### Variational inference: Strategy

- We'd like to compute  $p(\theta, z | x)$ , but it's too complicated.
- Strategy: Approximate as  $q(\theta, z)$  that has a "nice" form
- q is a function of variational parameters, optimized for each x.
- Maximize a lower bound on p(x).

The ELBO is the following lower bound on  $\log p(x)$ :

$$\log p(x) = \int \sum_{z} q(z, \theta) \log p(x) d\theta$$

The ELBO is the following lower bound on  $\log p(x)$ :

$$\log p(x) = \int \sum_{z} q(z, \theta) \log p(x) d\theta$$

$$= \sum_{z} \int q(z, \theta) \log \left( \frac{p(x, z, \theta) q(z, \theta)}{p(z, \theta \mid x) q(z, \theta)} \right) d\theta$$

The ELBO is the following lower bound on  $\log p(x)$ :

$$\log p(x) = \int \sum_{z} q(z,\theta) \log p(x) d\theta$$

$$= \sum_{z} \int q(z,\theta) \log \left( \frac{p(x,z,\theta) q(z,\theta)}{p(z,\theta \mid x) q(z,\theta)} \right) d\theta$$

$$= \sum_{z} \int q(z,\theta) \log \left( \frac{p(x,z,\theta)}{q(z,\theta)} \right) d\theta + \sum_{z} \int q(z,\theta) \log \left( \frac{q(z,\theta)}{p(z,\theta \mid x)} \right) d\theta$$

Ę

The ELBO is the following lower bound on  $\log p(x)$ :

$$\log p(x) = \int \sum_{z} q(z,\theta) \log p(x) d\theta$$

$$= \sum_{z} \int q(z,\theta) \log \left( \frac{p(x,z,\theta) q(z,\theta)}{p(z,\theta \mid x) q(z,\theta)} \right) d\theta$$

$$= \sum_{z} \int q(z,\theta) \log \left( \frac{p(x,z,\theta)}{q(z,\theta)} \right) d\theta + \sum_{z} \int q(z,\theta) \log \left( \frac{q(z,\theta)}{p(z,\theta \mid x)} \right) d\theta$$

$$\geq \sum_{z} \int q(z,\theta) \log \left( \frac{p(x,z,\theta)}{q(z,\theta)} \right) d\theta$$

Ę

The ELBO is the following lower bound on  $\log p(x)$ :

$$\log p(x) = \int \sum_{z} q(z,\theta) \log p(x) d\theta$$

$$= \sum_{z} \int q(z,\theta) \log \left( \frac{p(x,z,\theta) q(z,\theta)}{p(z,\theta \mid x) q(z,\theta)} \right) d\theta$$

$$= \sum_{z} \int q(z,\theta) \log \left( \frac{p(x,z,\theta)}{q(z,\theta)} \right) d\theta + \sum_{z} \int q(z,\theta) \log \left( \frac{q(z,\theta)}{p(z,\theta \mid x)} \right) d\theta$$

$$\geq \sum_{z} \int q(z,\theta) \log \left( \frac{p(x,z,\theta)}{q(z,\theta)} \right) d\theta$$

$$= H(q) + \mathbb{E}_{q}(\log p(x,z,\theta))$$

Ę

The ELBO is the following lower bound on  $\log p(x)$ :

$$\log p(x) = \int \sum_{z} q(z,\theta) \log p(x) d\theta$$

$$= \sum_{z} \int q(z,\theta) \log \left( \frac{p(x,z,\theta) q(z,\theta)}{p(z,\theta \mid x) q(z,\theta)} \right) d\theta$$

$$= \sum_{z} \int q(z,\theta) \log \left( \frac{p(x,z,\theta)}{q(z,\theta)} \right) d\theta + \sum_{z} \int q(z,\theta) \log \left( \frac{q(z,\theta)}{p(z,\theta \mid x)} \right) d\theta$$

$$\geq \sum_{z} \int q(z,\theta) \log \left( \frac{p(x,z,\theta)}{q(z,\theta)} \right) d\theta$$

$$= H(q) + \mathbb{E}_{q}(\log p(x,z,\theta))$$

We maximize this over the parameters of q



The inequality above uses concavity of the logarithm:

$$\log\left(\sum_{\alpha} w_{\alpha} x_{\alpha}\right) \geq \sum_{\alpha} w_{\alpha} \log x_{\alpha}$$

So, if  $q_{\alpha} \geq 0$  and  $p_{\alpha} \geq 0$  sum (or integrate) to one, then

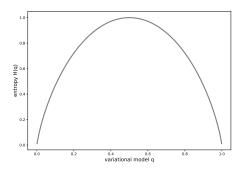
$$0 = \log \left( \sum_lpha oldsymbol{p}_lpha 
ight) = \log \left( \sum_lpha oldsymbol{q}_lpha rac{oldsymbol{p}_lpha}{oldsymbol{q}_lpha} 
ight) \geq \sum_lpha oldsymbol{q}_lpha \log \left( rac{oldsymbol{p}_lpha}{oldsymbol{q}_lpha} 
ight)$$

Therefore

$$D_{\mathit{KL}}(q \parallel p) \equiv \sum_{lpha} q_lpha \log \left(rac{q_lpha}{p_lpha}
ight) \geq 0$$

The ELBO is  $H(q) + \mathbb{E}_q(\log p)$ 

The entropy term H(q) encourages q to be spread out:



The cross-entropy  $\mathbb{E}_q \log p$  tries to match q to p

Since  $\log p(x, z, \theta) = \log p(z, \theta) + \log p(x \mid z, \theta)$ , the ELBO can be written as

$$\mathsf{ELBO} = \mathbb{E}_q(\log p(x \mid Z, \theta)) - D_{\mathsf{KL}}(q(Z, \theta) \parallel p(Z, \theta))$$

 The Kullback-Leibler divergence term acts as a regularizer, encouraging the variational distribution to be close to the prior

### **Example 2: A finite mixture model**

Fix two distributions  $F_0$  and  $F_1$ , with densities  $f_0(x)$  and  $f_1(x)$ , and form the mixture model

$$heta \sim \mathsf{Beta}(lpha,eta) \ X \, | \, heta \sim heta \mathsf{F}_1 + (1- heta) \mathsf{F}_0.$$

The likelihood for data  $x_1, \ldots, x_n$  is

$$p(x_{1:n}) = \int_0^1 \text{Beta}(\theta \mid \alpha, \beta) \prod_{i=1}^n (\theta f_1(x_i) + (1-\theta) f_0(x_i)) d\theta.$$

Our goal is to approximate the posterior  $p(\theta \mid x_{1:n})$ 

ć

### Variational approximation

Our variational approximation is

$$q(z,\theta) = q(\theta \mid \gamma_1, \gamma_2) \prod_{i=1}^n q_i^{z_i} (1-q_i)^{(1-z_i)}$$

where  $q(\theta \mid \gamma_1, \gamma_2)$  is a Beta $(\gamma_1, \gamma_2)$  distribution, and  $0 \le q_i \le 1$  are n free parameters.

Need to maximize ELBO  $H(q) + \mathbb{E}_q \log p$ 

### Variational algorithm for mixture

#### Variational inference

Iterate the following steps for variational parameters  $q_{1:n}$  and  $(\gamma_1, \gamma_2)$ :

**1** Holding  $q_i$  fixed, set  $\gamma = (\gamma_1, \gamma_2)$  to

$$\gamma_1 = \alpha + \sum_{i=1}^n q_i$$
  $\gamma_2 = \beta + n - \sum_{i=1}^n q_i$ 

2 Holding  $\gamma_1$  and  $\gamma_2$  fixed, set  $q_i$  to

$$q_i = \frac{f_1(x_i) \exp \psi(\gamma_1)}{f_1(x_i) \exp \psi(\gamma_1) + f_0(x_i) \exp \psi(\gamma_2)}$$

After convergence, approximate posterior distribution over  $\theta$  is

$$\widehat{p}(\theta \mid x_{1:n}) = \mathsf{Beta}(\theta \mid \gamma_1, \gamma_2)$$

### **Deterministic approximation**

- Convergence is numerical, not stochastic
- Posterior is approximated as a single Beta
- Very similar algorithm is used for topic models

# Some sample questions

Q: What is the best q we could use?

Q: What is the best *q* we could use?

A: The true posterior  $q(z, \theta \mid x) = p(z, \theta \mid x)$ 

Q: What is the best *q* we could use?

A: The true posterior  $q(z, \theta \mid x) = p(z, \theta \mid x)$ 

Why? Because this maximizes the ELBO. Mathematically,

$$\sum_{z} \int q(z,\theta) \log \left( \frac{q(z,\theta)}{p(z,\theta \mid x)} \right) d\theta = 0$$

in this case, so the ELBO inequality is an equality.

Q: How does the ELBO regularize?

Q: How does the ELBO regularize?

A: The entropy term favors distributions that are "spread out"

Q: How does the ELBO regularize?

A: The entropy term favors distributions that are "spread out"

Why? For discrete distributions the maximum entropy distribution is uniform. For Gaussian, the entropy is  $\log \sigma^2$  which favors  $\sigma^2$  large.

Q: Is the ELBO easy to maximize?

Q: Is the ELBO easy to maximize?

A: No.

Q: Is the ELBO easy to maximize?

A: No.

Why? In general it is non-convex, and the solution depends on where we start an iterative algorithm. This is unlike the Gibbs sampler, which converges to the right thing if we wait long enough.

#### Variational autoencoders

- Variational autoencoders are generative models that are trained using variational inference
- The "decoder" is a neural net that generates from a latent variable
- The "encoder" approximates the posterior distribution with another neural network trained using variational inference

### Variational autoencoders

Start with a generative model

$$z \sim N(0, I_k)$$
  
  $x \mid z \sim N(G(z), I_d)$ 

G(z) is the *generator network* or *decoder*. The latent z is k-dimensional and the output x is d-dimensional.

For example, use a 2-layer network

$$G(z) = A_2 \operatorname{ReLU}(A_1 z + b_1) + b_2$$

#### **Posterior inference**

How do we train the generative network?

$$Z \sim N(0, I_k)$$
  
 $X \mid z \sim N(G(z), I_d)$ 

We want the posterior distribution p(z | x)

But this is intractable, because  $G(\cdot)$  is nonlinear

Approach: Use variational inference

### **Using variational inference**

For variational inference we take

$$q(z \mid x) = N(\mu_x, \sigma_x^2 I_k)$$

where now  $\mu_x$  and  $\sigma_x^2$  are the *variational parameters* 

### **Using neural networks**

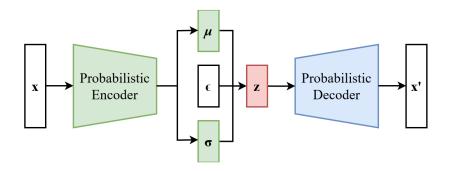
Rather than estimate  $\mu_X$  for each x, we build an encoder neural network that outputs the mean and variance.

For example:

$$\mu(x) = B_2 \text{ ReLU}(B_1 x + d_1) + d_2$$

and similarly for  $\log \sigma^2(x)$ .

### **VAE** architecture



## Training the neural networks

Decoder network:  $z \mapsto x$ , weights A, b

Encoder network:  $x \mapsto \mu(x)$ ,  $\log \sigma^2(x)$ , weights B, d

Train both networks simultaneously, to maximize the ELBO

- Decoder trained with  $\mathbb{E}_q \log p(x, Z_s)$  over weights A, b
- Encoder trained with  $H(q) + \mathbb{E}_q \log p(x, Z_s)$  over weights B, d

The entropy of a multivariate Gaussian with covariance  $\Sigma$  is  $\frac{1}{2} \log |\Sigma| + \text{constant}$ 

The entropy term here is  $\log \sigma^2(x) + \text{constant}$ , which is taken to be an output of the encoder network.

### Using neural networks

Now, approximate  $\mathbb{E}_q(\log p(x, Z))$  by sampling (weak law of large numbers)

$$\mathbb{E}_q(\log p(x,Z)) \approx \frac{1}{N} \sum_{s=1}^N \log p(x,Z_s)$$

Problem: The parameters of the recognition network have disappeared!

Solution: Reparameterize the samples by  $Z_i = \mu(x) + \sigma(x)\epsilon_i$  where  $\epsilon_i \sim N(0, I_k)$ .

This is called "the reparameterization trick." D. Kingma and M. Welling, "Autoencoding Variational Bayes," https://arxiv.org/abs/1312.6114.

## Simple example

Suppose  $x \mid z \sim N(G(z), I)$  where generator network is

$$G(z) = \text{ReLU}(Az + b).$$

Then 
$$-\log p(x \mid z)$$
 is

$$\frac{1}{2}||x - \text{ReLU}(Az + b)||^2 + \text{constant}$$

## Simple example

Suppose  $x \mid z \sim N(G(z), I)$  where generator network is

$$G(z) = \text{ReLU}(Az + b).$$

Then  $-\log p(x \mid z)$  is

$$\frac{1}{2}||x - \text{ReLU}(Az + b)||^2 + \text{constant}$$

And  $-\log p(z)$  is

$$\frac{1}{2}||z||^2 + \text{constant}$$

## Simple example (continued)

Next, suppose the approximate posterior is

$$q(z \mid x) = N(\mu(x), \sigma^2(x)I_k)$$

where encoder network is  $\mu(x) = \text{ReLU}(Bx + d)$ . Then

$$-\mathbb{E}_{q} \log p(x \mid Z) \approx \frac{1}{N} \sum_{s=1}^{N} \frac{1}{2} \|x - \text{ReLU}(AZ_{s} + b)\|^{2}$$

$$\stackrel{d}{=} \frac{1}{N} \sum_{s=1}^{N} \frac{1}{2} \|x - \text{ReLU}(A(\text{ReLU}(Bx + d) + \epsilon_{s}\sigma(x)) + \epsilon_{s}\sigma(x)) + \epsilon_{s}\sigma(x) + \epsilon_{s}\sigma$$

$$\stackrel{d}{=} \frac{1}{N} \sum_{s=1}^{N} \frac{1}{2} ||x - \text{ReLU}(A(\text{ReLU}(Bx + d) + \epsilon_s \sigma(x)) + b)||^2$$

### Simple example (continued)

Next, suppose the approximate posterior is

$$q(z \mid x) = N(\mu(x), \sigma^2(x)I_k)$$

where encoder network is  $\mu(x) = \text{ReLU}(Bx + d)$ . Then

$$-\mathbb{E}_q \log p(Z) = \frac{1}{2} \mathbb{E}_q ||Z||^2$$
$$= \frac{1}{2} ||\mu(x)||^2 + \frac{k}{2} \sigma^2(x)$$

### Resulting training objective

- This gives all of the pieces for the ELBO objective
- No need to work out the detailed calculations TensorFlow does this for you
- The ELBO is optimized by alternating stochastic gradient steps on the weights in the decoder and encoder networks

### Demo

#### Variational Autoencoder on MNIST data

This is a Tensorflow implementation of Variational Autoencoder (VAE) on MNIST data, based on Auto-Encoding Variational Bayes (Kingma and Welling 2014).

It uses probabilistic encoders and decoders realized by Multilayer Perceptrons (MLP) with a single hidden layer. The VAE was trained incrementally with minibatches using partial fit.

The MNIST dataset, distributed by Yann Lecun's THE MNIST DATARASE of handwritten digits website, consists of pair 'handwritten digit image' and 'label'. The image is a gray scale image with 28 x 28 pixels. Pixel values range from 0 (black) to 255 (white), scaled in the [0, 1] interval. The label is the actual digit, ranging from 10 x) the image represents.

The original notebook was composed by Jan Hendrik Metzen. Modifications were made by Sunnie Kim on May 24th, 2018.

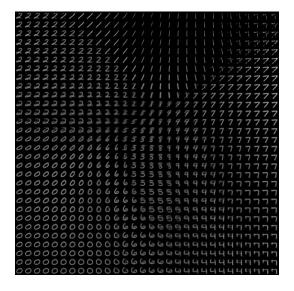
In (1): import numpy as np import tensorflow as ef import matploclib.pyplot as plt tmatploclib inline np.random.seed(0) f.set\_random.seed(0)

#### Load MNIST data in a format suited for Tensorflow

The 'input\_data' script is available at: https://raw.githubusercontent.com/tensorflow/tensorflow/master/tensorflow/examples/tutorials/mnist/input\_data.py

In [2]: from tensorflow.examples.tutorials.mnist import input\_data
tf.logging.set\_verbosity(tf.logging.ERROR)
mnist = input\_data.read\_data\_sets("NNIST\_data/", one\_hot=True)
n\_samples = mnist.train.num\_examples

### Visualizing a 2-dim latent space



#### **Deconvolutional neural networks**

- A convolutional network goes from a high dimensional input to a lower dimensional output
- The decoder / generator network goes from a low dimensional latent vector to a high dimensional output
- "Deconvolutional" or transposed convolutional networks are the "opposite" of CNNs
- Useful for this type of problem when data are images
- We won't dive into the details here

### **Summary**

- Variational methods make deterministic approximations
- General recipe: Maximize ELBO over variational parameters
- VAEs: Generator / decoder is a neural network
- Variational mean is output of a second neural network
- The two networks are trained together to maximize the ELBO