



S&DS 365 / 665
Intermediate Machine Learning

Discrete Data Graphs and Graph Neural Networks

October 26

Yale

A rare lull

- Assignment 3 out; due next Wednesday
- Assignment 4 posted next week



Graphs

- A natural language for describing various data
- Give information about relationships between variables
- Associated with each multivariate distribution

Undirected Graphs

A graph $G = (V, E)$ has vertices V , edges E .

If $X = (X_1, \dots, X_p)$ is a random variable, we will study graphs where there are p vertices, one for each X_j .

The graph will encode conditional independence relations among the variables.

Undirected graphs

Simplest case:



Here $V = \{X, Y, Z\}$ and $E = \{(X, Y), (Y, Z)\}$.

This encodes the independence relation

$$X \perp\!\!\!\perp Z \mid Y$$

which means that *X and Z are independent conditioned on Y*.

Markov Property


A probability distribution P satisfies the *global Markov property* with respect to a graph G if:

for any disjoint vertex subsets A , B , and C such that C separates A and B ,

$$X_A \perp\!\!\!\perp X_B \mid X_C.$$

- X_A are the random variables X_j with $j \in A$.
- C separates A and B means that there is no path from A to B that does not pass through C .

Example




$$X_1 \perp\!\!\!\perp X_4 \mid X_2, X_3$$


$$X_2 \perp\!\!\!\perp X_3 \mid X_1, X_4$$

Example: 2-dimensional grid

The blue node is independent of the red nodes given the white nodes.



Example



$C = \{3, 7\}$ separates $A = \{1, 2\}$ and $B = \{4, 8\}$. Hence,

$$\{X_1, X_2\} \perp\!\!\!\perp \{X_4, X_8\} \quad | \quad \{X_3, X_7\}$$

Special case

If $(i, j) \notin E$ then

$$X_i \perp\!\!\!\perp X_j \mid \{X_k : k \neq i, j\}$$

Special case

If $(i, j) \notin E$ then

$$X_i \perp\!\!\!\perp X_j \mid \{X_k : k \neq i, j\}$$

Lack of an edge from i to j implies that X_i and X_j are independent given all of the other random variables.

Graph estimation

- A graph G represents the class of distributions, $\mathcal{P}(G)$, the distributions that are Markov with respect to G
- Graph estimation: Given n samples $X_1, \dots, X_n \sim P$, estimate the graph G .

Factored form

Theorem (Hammersley, Clifford, Besag)

A positive distribution over random variables Z_1, \dots, Z_p satisfies the Markov properties of graph G if and only if it can be represented as

$$p(Z) \propto \prod_{c \in \mathcal{C}} \psi_c(Z_c)$$

where \mathcal{C} is the set of cliques in the graph G .

Gaussian case

Let $\Omega = \Sigma^{-1}$ be the precision matrix.

A zero in Ω indicates a *lack of the corresponding edge* in the graph

So, the adjacency matrix of the graph is

$$A = (\mathbb{1}(\Omega_{ij} \neq 0))$$

That is,

$$A_{ij} = \begin{cases} 1 & \text{if } |\Omega_{ij}| > 0 \\ 0 & \text{otherwise} \end{cases}$$


Gaussian case

$$\Omega \equiv \Sigma^{-1} = \begin{pmatrix} * & * & 0 \\ * & * & * \\ 0 & * & * \end{pmatrix}$$



Gaussian case

$$\Omega \equiv \Sigma^{-1} = \begin{pmatrix} * & * & * & 0 \\ * & * & 0 & * \\ * & 0 & * & * \\ 0 & * & * & * \end{pmatrix}$$



$$X_1 \perp\!\!\!\perp X_4 \mid X_2, X_3$$

Gaussian case: Algorithms

Two approaches:

- parallel lasso
- graphical lasso

Parallel Lasso:

- 1 For each $j = 1, \dots, p$ (in parallel): Regress X_j on all other variables using the lasso.
- 2 Put an edge between X_i and X_j if each appears in the regression of the other.

Graphical Lasso (glasso)

- Assume a multivariate Gaussian model
- Subtract out the sample mean
- Minimize the negative log-likelihood of the data, subject to a constraint on the sum of the absolute values of the inverse covariance

Graphical Lasso (glasso)

The glasso optimizes the parameters of $\Omega = \Sigma^{-1}$ by minimizing:

$$\text{trace}(\Omega S_n) - \log |\Omega| + \lambda \sum_{j \neq k} |\Omega_{jk}|$$

where $|\Omega|$ is the determinant and S_n is the sample covariance

$$S_n = \frac{1}{n} \sum_{i=1}^n x_i x_i^T$$

There is a blockwise gradient descent algorithm to minimize this, using iterative lassos



Discrete Graphical Models

Challenges of handling discrete data:

- Models don't have closed form; can't compute normalizing constant
- Need to use Gibbs sampling, variational inference
- No analogue of the graphical lasso

Discrete Graphical Models

- Positive distributions can be represented by an exponential family,

$$p(\mathbf{Z}; \beta) \propto \exp \left(\sum_{c \in \mathcal{C}} \beta_c \phi_c(\mathbf{Z}_c) \right)$$

- Special case: Ising Model (discrete Gaussian)

$$p(\mathbf{Z}; \beta) \propto \exp \left(\sum_{i \in V} \beta_i Z_i + \sum_{(i,j) \in E} \beta_{ij} Z_i Z_j \right).$$

Discrete Gaussian?

Note that we can write a multivariate Gaussian as follows:

$$p(\mathbf{z}) \propto \exp \left(\sum_i \beta_i z_i + \sum_{i,j} \beta_{ij} z_i z_j \right)$$

From edges to cliques

Take $\beta_i \equiv 0$ for simplicity

If we have a triangle (i, j, k) in the graph then the potential function corresponds to

$$\begin{aligned}\psi_{(ijk)}(Z_i, Z_j, Z_k) &= e^{\beta_{ij}Z_iZ_j} \cdot e^{\beta_{jk}Z_jZ_k} \cdot e^{\beta_{ik}Z_iZ_k} \\ &= e^{\beta_{ij}Z_iZ_j + \beta_{jk}Z_jZ_k + \beta_{ik}Z_iZ_k}\end{aligned}$$

Ising

We have a graph with edges E and vertices V . Each node i has a random variable Z_i that can be “up” ($Z_i = 1$) or “down” ($Z_i = 0$)

$$\mathbb{P}_\beta(z_1, \dots, z_n) \propto \exp \left(\sum_{s \in V} \beta_s z_s + \sum_{(s,t) \in E} \beta_{st} z_s z_t \right)$$

Since $2Z_i - 1 \in \{-1, 1\}$ if $Z_i \in \{0, 1\}$, can re-parameterize in terms of sample space $Z_i = \pm 1$.

Ising

We have a graph with edges E and vertices V . Each node i has a random variable Z_i that can be “up” ($Z_i = 1$) or “down” ($Z_i = 0$)

$$\mathbb{P}_{\beta}(z_1, \dots, z_n) \propto \exp \left(\sum_{s \in V} \beta_s z_s + \sum_{(s,t) \in E} \beta_{st} z_s z_t \right)$$

E are the set of edges, V are the vertices. Imagine the Z_i are votes of politicians, and the edges encode the social network of party affiliations

Since $2Z_i - 1 \in \{-1, 1\}$ if $Z_i \in \{0, 1\}$, can re-parameterize in terms of sample space $Z_i = \pm 1$.

Stochastic approximation

Gibbs sampler

Iterate until converged:

- 1 Choose vertex $s \in V$ at random
- 2 Sample z_s holding others fixed

$$\theta_s = \text{sigmoid} \left(\beta_s + \sum_{t \in N(s)} \beta_{st} z_t \right)$$
$$Z_s \mid \theta_s \sim \text{Bernoulli}(\theta_s)$$

Deterministic approximation

Mean field variational algorithm


Iterate until converged:

- 1 Choose vertex $s \in V$ at random
- 2 Update mean μ_s holding others fixed


$$\mu_s = \text{sigmoid} \left(\beta_s + \sum_{t \in N(s)} \beta_{st} \mu_t \right)$$

Graph Estimation


- Given n i.i.d. samples from an Ising distribution, $\{Z_i, i = 1, \dots, n\}$, (each is a p -vector of $\{0, 1\}$ values) identify underlying graph



- Multiple examples are observed:



Local Distributions



- Consider Ising model $p_{\beta}(Z) \propto \exp \left(\sum_{i \in V} \beta_i Z_i + \sum_{(i,j) \in E} \beta_{ij} Z_i Z_j \right)$.
- Conditioned on (z_2, \dots, z_p) , variable $Z_1 \in \{0, 1\}$ has probability mass function given by a logistic function,


$$\mathbb{P}(Z_1 = 1 \mid z_2, \dots, z_p) = \text{sigmoid} \left(\beta_1 + \sum_{j \in \mathcal{N}(1)} \beta_{1j} z_j \right)$$

Parallel lasso (sparse logistic regressions)

Strategy

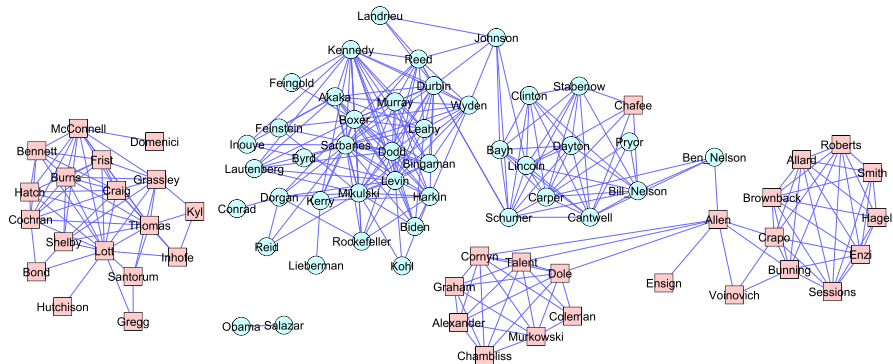
- Perform ℓ_1 regularized logistic regression of each node Z_i on $Z_{\setminus i} = \{Z_j, j \neq i\}$ to estimate neighbors $\hat{\mathcal{N}}(i)$
- Two versions:
 - ▶ Create an edge (i, j) if $j \in \hat{\mathcal{N}}(i)$ **and** $i \in \hat{\mathcal{N}}(j)$
 - ▶ Create an edge (i, j) if $j \in \hat{\mathcal{N}}(i)$ **or** $i \in \hat{\mathcal{N}}(j)$

S&P 500: Ising Model (Price up or down?)



Voting Data

Voting records of US Senate, 2006-2008



Scaling behavior: Performance with data size

Maximum degree d of the p variables. Sample size n must satisfy

$$\text{Ising model: } n \geq d^3 \log p$$

$$\text{Graphical lasso: } n \geq d^2 \log p$$

$$\text{Parallel lasso: } n \geq d \log p$$

$$\text{Lower bound: } n \geq d \log p$$

- Each method makes different *incoherence assumptions*:
 - ▶ Correlations between unrelated variables not too large

Summary: Graphs for discrete data

- A positive distribution factors into product of potential functions on the cliques of the graph
- Graphs and independence relations are same for discrete data
- Ising models are discrete Gaussians
- No version of the graphical lasso holds for discrete data; instead, we use the parallel lasso

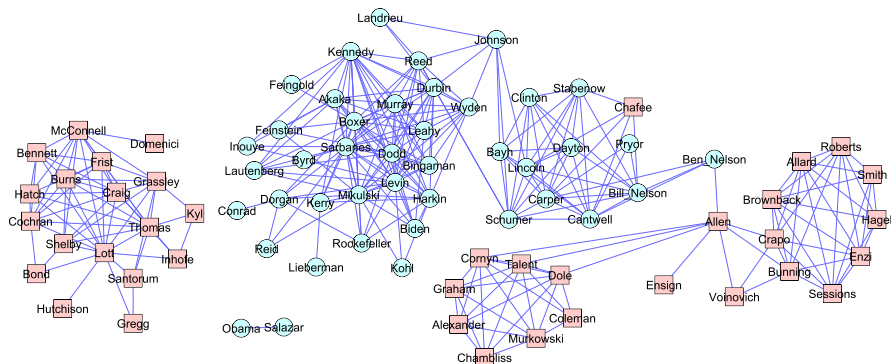
Graph neural networks

Next, we'll discuss graph neural networks, following this article:

<https://distill.pub/2021/understanding-gnns/>

Recall: Voting Data

Voting records of US Senate, 2006-2008




Suppose we have the graph and other covariates for each node or edge. How can we classify the senators according to political orientation?



Graph neural networks

- We'll discuss how GNNs correspond to CNNs
- The graph Laplacian plays a central role


Equivariance problem



Equivariance problem



Graph Laplacian



$$\begin{array}{c} \text{A} \\ \text{B} \\ \text{C} \\ \text{D} \\ \text{E} \\ \text{F} \\ \text{G} \end{array} \begin{bmatrix} & \text{A} & \text{B} & \text{C} & \text{D} & \text{E} & \text{F} & \text{G} \\ \begin{bmatrix} 1 & -1 & & & & & \\ -1 & 2 & -1 & & & & \\ & -1 & 5 & -1 & -1 & -1 & -1 \\ & & -1 & 1 & & & \\ & & -1 & & 1 & & \\ & & -1 & & & 1 & \\ & & -1 & & & & 1 \end{bmatrix} \end{bmatrix}$$

Laplacian L of G

Polynomials of the Laplacian

$$p_w(L) = w_0 I_n + w_1 L + w_2 L^2 + \cdots w_d L^d$$

If $\text{dist}(u, v) > i$ then the (u, v) entry of L^i is zero


- This is analogous to a CNN filter (kernel)
- The weights w_i play role of filter coefficients
- Degree d of polynomial plays role of the size of the kernel

The Laplacian is a Mercer kernel

- Symmetric $L_{uv} = L_{vu}$
- Positive-definite:


$$f^T L f = \sum_{(u,v) \in E} (f_u - f_v)^2 \geq 0$$

Sanity checks



What is the Laplacian L ?


Sanity checks



What is the Laplacian L ?


$$L = \begin{pmatrix} 2 & -1 & -1 & 0 \\ -1 & 2 & 0 & -1 \\ -1 & 0 & 2 & -1 \\ 0 & -1 & -1 & 2 \end{pmatrix}$$

Sanity checks



What is L^2 ?


Sanity checks



What is L^2 ?


$$L^2 = \begin{pmatrix} 6 & -4 & -4 & 2 \\ -4 & 6 & 2 & -4 \\ -4 & 2 & 6 & -4 \\ 2 & -4 & -4 & 6 \end{pmatrix}$$

Sanity checks



If $x = (1, 2, 3, 4)^T$ what is $h = \text{ReLU}(Lx)$?


Sanity checks



If $x = (1, 2, 3, 4)^T$ what is $h = \text{ReLU}(Lx)$?


$$\text{ReLU}(Lx) = \text{ReLU}((-3, -1, 1, 3)^T) = (0, 0, 1, 3)^T$$

Sanity checks



If $x = (1, 2, 3, 4)^T$ what is $x^T L x$?

Sanity checks



If $x = (1, 2, 3, 4)^T$ what is $x^T L x$?

$$x^T L x = \sum_{(u,v) \in E} (x_u - x_v)^2 = 10$$

Whence equivariance

A transformation $f : \mathbb{R}^n \longrightarrow \mathbb{R}^n$ is equivariant if

$$f(Px) = Pf(x)$$

for any permutation matrix P , where $PP^T = I$.

The transformed data and Laplacian are

$$x \longrightarrow Px$$

$$L \longrightarrow PLP^T$$

$$L^i \longrightarrow PL^iP^T$$

Whence equivariance

A transformation $f : \mathbb{R}^n \rightarrow \mathbb{R}^n$ is equivariant if

$$f(Px) = Pf(x)$$

for any permutation matrix P , where $PP^T = I$.

The transformed polynomial kernels are

$$\begin{aligned} f(Px) &= \sum_{i=0}^d w_i (PL^i P^T) Px \\ &= \sum_{i=0}^d w_i PL^i x \\ &= P \sum_{i=0}^d w_i L^i x \\ &= Pf(x) \end{aligned}$$

Building layers

Let $h^{(k)}$ be the neurons at layer k .


We start with $h^{(0)} = x$, a value x_j at each node j

The next layer is

$$h^{(k+1)} = \varphi \left(p_w(L) h^{(k)} \right)$$

See tutorial for other ways of building layers

Building layers



Course ad

CPSC 483: Deep Learning on Graphs

Instructor: Rex Ying

Summary: Graph neural nets

- Certain data have natural graphical structure
- GNNs are analogues of CNNs for graphs
- Based on use of graph Laplacian
- Independent of ordering of nodes (equivariant)

Next topic: Reinforcement learning