

S&DS 365 / 665  
Intermediate Machine Learning

# Smoothing and Density Estimation

September 12

Yale

# Topics for today

- Recap: Smoothing kernels
- Kernel density estimation
- Bias-variance decomposition
- Intro to Mercer kernels

# Some reminders

- Quiz 1: Great job!
- Assn 1 posted on Wednesday
- Topics: Lasso, smoothing, Mercer kernels, some neural nets
- Questions?

# Notes

- Notes posted to course page  
`http://interml.ydata123.org`
- Readings from “Probabilistic Machine Learning: An Introduction”
- `https://probml.github.io/pml-book/book1.html`

# Nonparametric Regression

Given  $(X_1, Y_1), \dots, (X_n, Y_n)$  predict  $Y$  from  $X$ .

Assume only that  $Y_i = m(X_i) + \epsilon_i$  where  $m(x)$  is a smooth function of  $x$ .

The most popular methods are *kernel methods*. However, there are two types of kernels:

- 1 Smoothing kernels
- 2 Mercer kernels

Smoothing kernels involve local averaging.  
Mercer kernels involve regularization.

# Smoothing Kernels

- Smoothing kernel estimator:

$$\hat{m}_h(x) = \frac{\sum_{i=1}^n Y_i K_h(X_i, x)}{\sum_{i=1}^n K_h(X_i, x)}$$

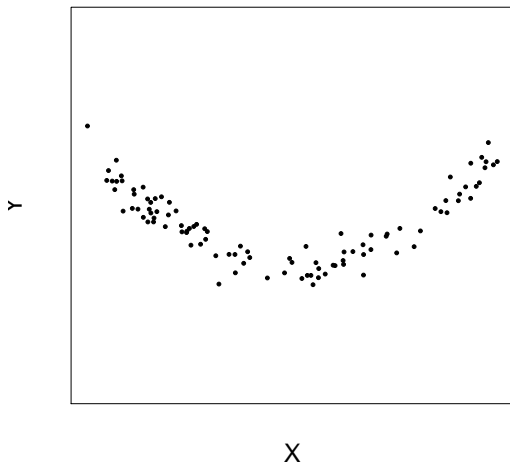
where  $K_h(x, z)$  is a *kernel* such as

$$K_h(x, z) = \exp\left(-\frac{\|x - z\|^2}{2h^2}\right)$$


and  $h > 0$  is called the *bandwidth*.

- $\hat{m}_h(x)$  is just a local average of the  $Y_i$ 's near  $x$ .
- The bandwidth  $h$  controls the bias-variance tradeoff:  
*Small  $h$  = large variance* while *large  $h$  = large bias*.

## Example: Some Data – Plot of $Y_i$ versus $X_i$

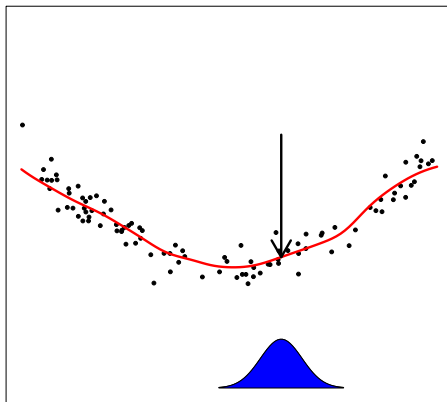


**Example:**  $\hat{m}(x)$







$\hat{m}(x)$  is a local average




# Effect of the bandwidth $h$




very small bandwidth



small bandwidth



medium bandwidth



large bandwidth

**Let's revisit the notebook**

# Smoothing Kernels

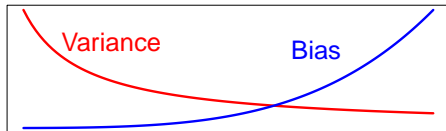
$$\text{Risk} = \mathbb{E}(Y - \hat{m}_h(X))^2 = \text{bias}^2 + \text{variance} + \sigma^2.$$

$\sigma^2 = \mathbb{E}(Y - m(X))^2$  is the unavoidable prediction error.

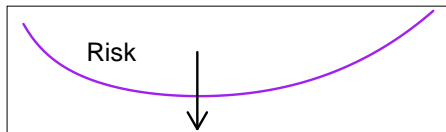
*small h*: low bias, high variance (undersmoothing)

*large h*: high bias, low variance (oversmoothing)

# Risk Versus Bandwidth



$h$



optimal  $h$

# Estimating the Risk: Cross-Validation

To choose  $h$  we need to estimate the risk  $R(h)$ . We can estimate the risk by using *cross-validation*.

- 1 Omit  $(X_i, Y_i)$  to get  $\hat{m}_{h,(i)}$ , then predict:  $\hat{Y}_{(i)} = \hat{m}_{h,(i)}(X_i)$ .
- 2 Repeat this for all observations.
- 3 The cross-validation estimate of risk is:

$$\hat{R}(h) = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_{(i)})^2.$$

*Shortcut formula:* Whenever  $\hat{Y} = LY$  we can use the shortcut

$$\hat{R}(h) = \frac{1}{n} \sum_{i=1}^n \left( \frac{Y_i - \hat{Y}_i}{1 - L_{ii}} \right)^2.$$

In this case  $L_{ij} = K_h(X_i, X_i) / \sum_t K_h(X_i, X_t)$ .

# Shortcut formula

Let's prove the shortcut formula. Let  $K_{ij} = K_h(X_i, X_j)$ . We have

$$\begin{aligned}\hat{Y}_{(i)} &= \frac{\sum_{j \neq i} K_{ij} Y_j}{\sum_{j \neq i} K_{ij}} \\&= \frac{\sum_j K_{ij} Y_j - K_{ii} Y_i}{\sum_j K_{ij} - K_{ii}} \\&= \frac{\sum_j L_{ij} Y_j - L_{ii} Y_i}{1 - L_{ii}} \\&= \frac{\hat{Y}_i - L_{ii} Y_i}{1 - L_{ii}}\end{aligned}$$

# Shortcut formula

It follows that


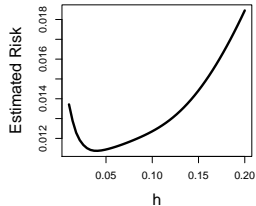
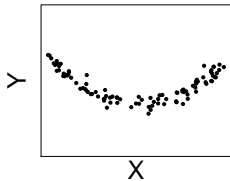
$$\begin{aligned}\left(Y_i - \hat{Y}_{(i)}\right)^2 &= \left(Y_i - \frac{\hat{Y}_i - L_{ii} Y_i}{1 - L_{ii}}\right)^2 \\ &= \left(\frac{Y_i - \hat{Y}_i}{1 - L_{ii}}\right)^2\end{aligned}$$



# Summary so far

- 1 Compute  $\hat{m}_h$  for each  $h$
- 2 Estimate the risk  $\hat{R}(h)$  using LOOCV
- 3 Choose bandwidth  $\hat{h}$  to minimize  $\hat{R}(h)$
- 4 Let  $\hat{m}(x) = \hat{m}_{\hat{h}}(x)$

# Example



# Kernel density estimation

To estimate a density, use the same idea behind kernel smoothing:

$$\begin{aligned}\hat{f}(x) &= \frac{1}{n} \sum_{i=1}^n K_h(X_i, x) \\ &= \frac{1}{n} \sum_{i=1}^n \frac{1}{h} K\left(\frac{X_i - x}{h}\right)\end{aligned}$$

We require that  $\int K(u) du = 1$  and  $K \geq 0$  is symmetric around zero (an even function).

This places a “bump function” around each data point, and averages them (a mixture model)

# Kernel density estimation

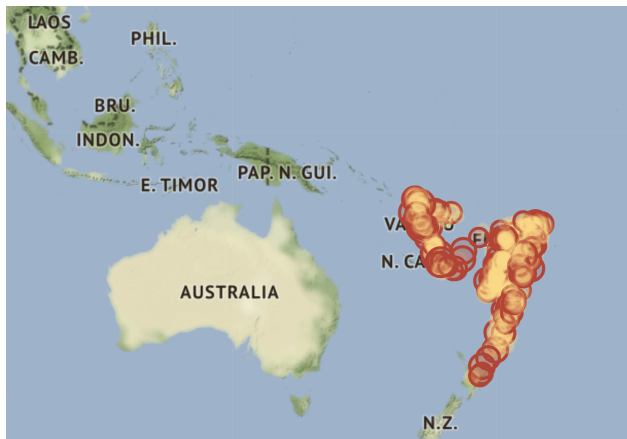
In  $p$  dimensions:

$$\begin{aligned}\hat{f}(x) &= \frac{1}{n} \sum_{i=1}^n K_h(X_i, x) \\ &= \frac{1}{n h^p} \sum_{i=1}^n K\left(\frac{X_i - x}{h}\right)\end{aligned}$$

We require that  $\int K(u) du = 1$  and  $K$  is symmetric around zero.

This places a “bump function” around each data point, and averages them (a mixture model)

# KDE demo: Fiji earthquakes



# Kernel density estimation

The bias-variance tradeoff:

$$\text{bias}^2(x) \approx h^4$$

$$\text{var}(x) \approx \frac{1}{n h^p}$$

Note that the variance scales according to the expected number of data points in a cube of side length  $h$  in  $p$ -dimensions.

We'll go through the calculation of this on the board. Notes are posted to <http://interml.ydata123.org>

# Back to regression

Using a kernel density estimator, the “plug-in” regression estimate gives us back the kernel smoother:

$$\begin{aligned}\hat{m}(x) &= \int y \hat{f}(y | x) dy \\ &= \frac{\int y \hat{f}(x, y) dy}{\hat{f}(x)} \\ &= \frac{\sum_i Y_i K_h(X_i, x)}{\sum_i K_h(X_i, x)}\end{aligned}$$

# Generative models

- A density estimate is a *generative model*
- We can sample from the density to “generate” a new data point
- What is an algorithm for sampling from the estimated distribution?



# Generative models

- 1 Sample an index  $i$  uniformly from 1 to  $n$
- 2 Sample a point  $x$  from a Gaussian with mean  $X_i$  and variance  $h^2$

# Generative models

Some recent generative models for images:



# Generative models

Some recent generative models for images:



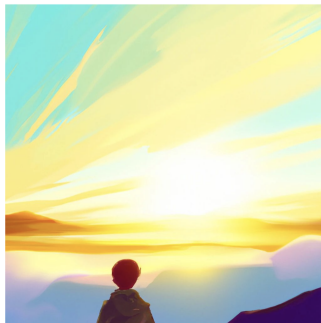
[API](#) [RESEARCH](#) [BLOG](#) [ABOUT](#)

## DALL·E Now Available in Beta

We'll invite 1 million people from our waitlist over the coming weeks. Users can create with DALL·E using free credits that refill every month, and buy additional credits in 115-generation increments for \$15.

[JOIN DALL·E 2 WAITLIST](#)

July 20, 2022  
5 minute read



# Summary

- Smoothing methods compute local averages, weighting points by a kernel
- Shape of the kernel doesn't matter (much)
- KDE places a density around each data point, and averages
- The COD limits use of both approaches to low dimensions